

Beyond Classical Search

CHAPTER 4
COSC 370
SPRING 2019
ALAN C. JAMIESON
SOME SLIDE CONTENT FROM RUSSELL &
NORVIG PROVIDED SLIDES

1

- Hill-climbing
- Simulated annealing
- Local beam search
- Genetic algorithms
- Searching with non-determinism
- Searching with partial observations

2


Local Searches

- Sometimes, we don't need to examine the whole state space (e.g. TSP).
- Iterative improvement algorithms – class of algorithms that focus on getting to the end state, rather than the path.
- State space = all possible complete configurations
- Optimize over those configurations!
- Constant space requirements.

3

Example: TSP

- Start with complete tour, perform pairwise exchange:




- Variants of this approach get near optimal solutions very quickly with thousands of cities.

4

Example: n-Queens

- Place n queens on an $n \times n$ board with no queen able to attack another.
- Strategy here?
 - Move a single queen to reduce conflicts:



- Very quick solution!

5

Hill-Climbing


- a.k.a. Gradient Ascent/Descent

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```

6

Watch the Landscape!

- Consider that the state space has a location and an elevation.
- Location = state itself
- Elevation = evaluation function result




7

Local Extremal States

- Local maximal/minimal states = bad news.

18	12	14	13	12	14	14
14	16	13	15	12	14	16
14	12	18	13	15	12	14
15	14	14	13	15	16	13
17	14	17	15	16	14	16
17	16	16	18	15	15	16
18	14	16	15	14	16	16
14	14	13	17	16	14	18



- How to get around? Random restart & random move.

8

Simulated Annealing

- Notice – hill-climbing techniques that never make a downhill move can be incomplete.
- Annealing – process used to temper or harden metals by heating them to high temperatures then gradually cooling them.
- Simulated annealing – gradient descent algorithm that incorporates some random moves to keep an agent out of local minimums/maximums.
- Allow “bad” moves but gradually decrease their frequency and size.

9

Simulated Annealing

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
inputs: problem, a problem
       schedule, a mapping from time to “temperature”
local variables: current, a node
                next, a node
                T, a “temperature” controlling prob. of downward steps

current ← MAKE-NODE(INITIAL-STATE[problem])
for t ← 1 to ∞ do
  T ← schedule[t]
  if T = 0 then return current
  next ← a randomly selected successor of current
   $\Delta E \leftarrow \text{VALUE}[\text{next}] - \text{VALUE}[\text{current}]$ 
  if  $\Delta E > 0$  then current ← next
  else current ← next only with probability  $e^{\Delta E/T}$ 

```

10

Local Beam Search (briefly)

- Simulated Annealing = one state kept in memory.
- What if we were able to have k states? Would this speed things up?
- Local beam search – start with k randomly generated initial states and expand. Choose the k best evaluated nodes expanded and continue until you get the goal.
- Stochastic beam search expands this further by utilizing a probability to choose k random expanded states.


11

Genetic Algorithms (briefly)

- A expansion of the stochastic beam search.
- Choose population (randomly selected states).
- Each individual state must be represented by a string over a finite alphabet (binary bits, characters, etc.)
 - For instance: n -queens states represent by bits the location of each queen, 8-queens needs 24 bits, 3 bits per column.
- The next generation of states is created by examining the current set of states and their fitness (typically by the standard evaluative function). Then a random (but probabilistically chosen) set of states are chosen. Those states are paired and a crossover point is chosen randomly. New states are created from the two parents by appending substrings from the parents together based on the crossover.

12

Genetic Algorithms (briefly)



- Mutations can be added in with small independent probability. This will be dependent on the problem, but represent a change in a generated state's string. For instance – moving a single queen randomly in its column.

13


Adding Nondeterminism

- Consider an erratic Vacuum World - when the *Suck* command is executed the vacuum cleans just the current square, cleans the current square and an adjacent square, or deposits dirt onto a clean square.
- How do we deal with this non-determinism? What does our state space look like? How do we search when we don't know what's next?
- Consider all the options!

14

AND-OR Search Trees

- Search trees where certain branches represent decisions (OR) and certain branches represent options that must all be taken.
- For instance: our *Suck* command on a clean square.
- What about just sitting at a particular state?
- We assume that we have a fully observable environment.



15


16

What about...

- A slippery Vacuum World? Movement actions occasionally fail.
- Try, try again – an extension of our *if-else* state based operations before, occasionally we will have to loop back to the state where the failure occurred, then continue to execute as normal.
- Requires some knowledge of whether or not an action failed.



- How do we deal with an agent that doesn't provide enough information to set state?
- Belief states: a set of fully observable states representing all possible configurations of state.
- Predictions if we have partial or no sensor ability!



17

18

Quick Exercise!

- 4.1 a-d

19